

Data 608 Module 2

September 30, 2019

```
[5]: import plotly
```

```
[6]: import datashader as ds
import datashader.transfer_functions as tf
import datashader.glyphs
from datashader import reductions
from datashader.core import bypixel
from datashader.utils import lnglat_to_meters as webm, export_image
from datashader.colors import colormap_select, Greys9, viridis, inferno
import copy

from pyproj import Proj, transform
import numpy as np
import pandas as pd
import urllib
import json
import datetime
import colorlover as cl

import plotly.offline as py
import plotly.graph_objs as go
from plotly import tools

from shapely.geometry import Point, Polygon, shape
# In order to get shapely, you'll need to run [pip install shapely.geometry]
↳ from your terminal

from functools import partial

from IPython.display import GeoJSON
```

```
[45]: # Code to read in v17, column names have been updated (without upper case
↳ letters) for v18

bk = pd.read_csv('/Users/christinakasman/Desktop/DATA 608 Pluto/BK2017V11.csv')
bx = pd.read_csv('/Users/christinakasman/Desktop/DATA 608 Pluto/BX2017V11.csv')
```

```

mn = pd.read_csv('/Users/christinakasman/Desktop/DATA 608 Pluto/MN2017V11.csv')
qn = pd.read_csv('/Users/christinakasman/Desktop/DATA 608 Pluto/QN2017V11.csv')
si = pd.read_csv('/Users/christinakasman/Desktop/DATA 608 Pluto/SI2017V11.csv')

ny = pd.concat([bk, bx, mn, qn, si], ignore_index=True)

#ny = pd.read_csv('nyc_pluto_18v2_csv/pluto_18v2.csv')

# Getting rid of some outliers
ny = ny[(ny['YearBuilt'] > 1850) & (ny['YearBuilt'] < 2020) & (ny['NumFloors'] !
↪= 0)]

ny.head()

```

/Users/christinakasman/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: FutureWarning:

Sorting because non-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

```

[45]:
APPBBL APPDate Address AreaSource AssessLand AssessTot \
1 0.0 NaN 10 JAY STREET 2 834300 10156950
17 0.0 NaN 1 PLYMOUTH STREET 2 129600 351000
20 0.0 NaN 65 PLYMOUTH STREET 2 178200 236700
22 0.0 NaN 135 PLYMOUTH STREET 2 514350 7785450
23 0.0 NaN 20 JAY STREET 2 972000 40980600

BBL BldgArea BldgClass BldgDepth ... YearAlter1 YearAlter2 \
1 3000010050 163894 06 195.0 ... 1994 2015
17 3000070021 9585 G1 113.0 ... 1988 0
20 3000160005 5000 Q0 165.0 ... 0 0
22 3000180001 211386 D5 48.0 ... 2014 0
23 3000190001 500000 06 207.0 ... 1999 0

YearBuilt ZMCode ZipCode ZoneDist1 ZoneDist2 ZoneDist3 ZoneDist4 \
1 1920 NaN 11201.0 M1-4/R8A M3-1 NaN NaN
17 1988 NaN 11201.0 M3-1 NaN NaN NaN
20 1920 NaN 11201.0 PARK NaN NaN NaN
22 1900 NaN 11201.0 M1-4/R8A NaN NaN NaN
23 1911 NaN 11201.0 M1-4/R8A NaN NaN NaN

```

	ZoneMap
1	12d
17	12d
20	12d
22	12d
23	12d

[5 rows x 89 columns]

```
[13]: wgs84 = Proj("+proj=longlat +ellps=GRS80 +datum=NAD83 +no_defs")
nyli = Proj("+proj=lcc +lat_1=40.66666666666666 +lat_2=41.03333333333333
↳+lat_0=40.16666666666666 +lon_0=-74 +x_0=300000 +y_0=0 +ellps=GRS80
↳+datum=NAD83 +to_meter=0.3048006096012192 +no_defs")
ny['XCoord'] = 0.3048*ny['XCoord']
ny['YCoord'] = 0.3048*ny['YCoord']
ny['lon'], ny['lat'] = transform(nyli, wgs84, ny['XCoord'].values, ny['YCoord'].
↳values)

ny = ny[(ny['lon'] < -60) & (ny['lon'] > -100) & (ny['lat'] < 60) & (ny['lat']
↳> 20)]

#Defining some helper functions for DataShader
background = "black"
export = partial(export_image, background = background, export_path="export")
cm = partial(colormap_select, reverse=(background!="black"))
```

Part 1: Binning and Aggregation Binning is a common strategy for visualizing large datasets. Binning is inherent to a few types of visualizations, such as histograms and 2D histograms (also check out their close relatives: 2D density plots and the more general form: heatmaps).

While these visualization types explicitly include binning, any type of visualization used with aggregated data can be looked at in the same way. For example, lets say we wanted to look at building construction over time. This would be best viewed as a line graph, but we can still think of our results as being binned by year:

```
[14]: trace = go.Scatter(
    # I'm choosing BBL here because I know it's a unique key.
    x = ny.groupby('YearBuilt').count()['BBL'].index,
    y = ny.groupby('YearBuilt').count()['BBL']
)

layout = go.Layout(
    xaxis = dict(title = 'Year Built'),
    yaxis = dict(title = 'Number of Lots Built')
)

fig = go.Figure(data = [trace], layout = layout)
```

```
py.iplot(fig)
```

Something looks off... You're going to have to deal with this imperfect data to answer this first question.

But first: some notes on pandas. Pandas dataframes are a different beast than R dataframes, here are some tips to help you get up to speed:

Hello all, here are some pandas tips to help you guys through this homework:

Indexing and Selecting: `.loc` and `.iloc` are the analogs for base R subsetting, or `filter()` in dplyr

Group By: This is the pandas analog to `group_by()` and the appended function the analog to `summarize()`. Try out a few examples of this, and display the results in Jupyter. Take note of what's happening to the indexes, you'll notice that they'll become hierarchical. I personally find this more of a burden than a help, and this sort of hierarchical indexing leads to a fundamentally different experience compared to R dataframes. Once you perform an aggregation, try running the resulting hierarchical dataframe through a `reset_index()`.

Reset_index: I personally find the hierarchical indexes more of a burden than a help, and this sort of hierarchical indexing leads to a fundamentally different experience compared to R dataframes. `reset_index()` is a way of restoring a dataframe to a flatter index style. Grouping is where you'll notice it the most, but it's also useful when you filter data, and in a few other split-apply-combine workflows. With pandas indexes are more meaningful, so use this if you start getting unexpected results.

Indexes are more important in Pandas than in R. If you delve deeper into the using python for data science, you'll begin to see the benefits in many places (despite the personal gripes I highlighted above.) One place these indexes come in handy is with time series data. The pandas docs have a huge section on datetime indexing. In particular, check out `resample`, which provides time series specific aggregation.

Merging, joining, and concatenation: There's some overlap between these different types of merges, so use this as your guide. `Concat` is a single function that replaces `cbind` and `rbind` in R, and the results are driven by the indexes. Read through these examples to get a feel on how these are performed, but you will have to manage your indexes when you're using these functions. Merges are fairly similar to merges in R, similarly mapping to SQL joins.

Apply: This is explained in the "group by" section linked above. These are your analogs to the `plyr` library in R. Take note of the lambda syntax used here, these are anonymous functions in python. Rather than predefining a custom function, you can just define it inline using `lambda`.

Browse through the other sections for some other specifics, in particular reshaping and categorical data (pandas' answer to factors.) Pandas can take a while to get used to, but it is a pretty strong framework that makes more advanced functions easier once you get used to it. Rolling functions for example follow logically from the apply workflow (and led to the best google results ever when I first tried to find this out and googled "pandas rolling")

Google Wes McKinney's book "Python for Data Analysis," which is a cookbook style intro to pandas. It's an O'Reilly book that should be pretty available out there.

Question After a few building collapses, the City of New York is going to begin investigating older buildings for safety. The city is particularly worried about buildings that were unusually tall when they were built, since best-practices for safety hadn't yet been determined. Create a graph that shows how many buildings of a certain number of floors were built in each year (note: you may want to use a log scale for the number of buildings). Find a strategy to bin buildings (It should be clear 20-29-story buildings, 30-39-story buildings, and 40-49-story buildings were first built in large numbers, but does it make sense to continue in this way as you get taller?)

```
[113]: ny['XCoord'].describe()
```

```
[113]: count      8.117820e+05
      mean       1.006355e+06
      std        3.243865e+04
      min        9.131230e+05
      25%        9.895860e+05
      50%        1.008949e+06
      75%        1.029431e+06
      max        1.067279e+06
      Name: XCoord, dtype: float64
```

```
[114]: ny['YCoord'].describe()
```

```
[114]: count      811782.000000
      mean      191449.447992
      std       30465.111991
      min      120308.000000
      25%      168169.000000
      50%      189198.000000
      75%      210843.000000
      max      272275.000000
      Name: YCoord, dtype: float64
```

```
[50]: grouped = ny.groupby(['YearBuilt', 'NumFloors']).count()
      grouped.head()
```

```
[50]:
```

		APPBBL	APPDate	Address	AreaSource	AssessLand	\
YearBuilt	NumFloors						
1851	3.0	3	0	3	3	3	
	4.0	4	1	4	4	4	
	5.0	1	1	1	1	1	
1852	2.0	2	0	2	2	2	
	3.0	7	0	7	7	7	

		AssessTot	BBL	BldgArea	BldgClass	BldgDepth	...	\
YearBuilt	NumFloors						...	
1851	3.0	3	3	3	3	3	...	
	4.0	4	4	4	4	4	...	

	5.0	1	1	1	1	1	...
1852	2.0	2	2	2	2	2	...
	3.0	7	7	7	7	7	...

		YCoord	YearAlter1	YearAlter2	ZMCode	ZipCode	\
YearBuilt	NumFloors						
1851	3.0	3	3	3	0	3	
	4.0	4	4	4	0	4	
	5.0	1	1	1	0	1	
1852	2.0	2	2	2	0	2	
	3.0	7	7	7	0	7	

		ZoneDist1	ZoneDist2	ZoneDist3	ZoneDist4	ZoneMap
YearBuilt	NumFloors					
1851	3.0	3	0	0	0	3
	4.0	4	0	0	0	4
	5.0	1	0	0	0	1
1852	2.0	2	0	0	0	2
	3.0	7	0	0	0	7

[5 rows x 87 columns]

```
[170]: ny2 = ny[['YearBuilt', 'NumFloors', 'BBL']].copy()
```

```
[174]: ny2['decadebuilt'] = (np.floor(ny_floors['YearBuilt'] / 10.0).astype(int) * 10)
ny2.head()
```

```
[174]:
```

	YearBuilt	NumFloors	BBL	decadebuilt
1	1920	10.0	3000010050	1920
17	1988	1.0	3000070021	1980
20	1920	1.0	3000160005	1920
22	1900	7.0	3000180001	1900
23	1911	11.0	3000190001	1910

```
[191]: ny2['NumFloors'].describe()
```

```
[191]:
```

count	814035.000000
mean	2.442139
std	1.936177
min	0.500000
25%	2.000000
50%	2.000000
75%	2.500000
max	119.000000

Name: NumFloors, dtype: float64

```
[189]: fig = px.box(ny2, x="decadebuilt", y="NumFloors")
fig.show()
```

```
[203]: ny2.sort_values(by=['YearBuilt'], inplace = True)
```

```
#Bin levels
Levels = {
    1: "< 5 Floors",
    2: "5-10 Floors",
    3: "10-15 Floors",
    4: "15-20 Floors",
    5: "> 20 Floors"

}

def floorLevel(floor):
    if floor <=5:
        return Levels[1]
    elif floor > 5 and floor <= 10:
        return Levels[2]
    elif floor > 10 and floor <= 15:
        return Levels[3]
    elif floor > 15 and floor <= 20:
        return Levels[4]
    elif floor > 20:
        return Levels[5]
```

```
[213]: ny2['level'] = ""
ny2['level'] = ny2['NumFloors'].apply(floorLevel)

ny2 = ny2[ny2.level != '< 5 Floors']
ny2 = ny2[ny2.level != '5-10 Floors']
ny2.head()
```

```
[213]:
```

	YearBuilt	NumFloors	BBL	decadebuilt	level
408723	1853	11.0	1002177501	1850	10-15 Floors
388488	1868	12.0	1013940023	1860	10-15 Floors
70164	1885	11.0	3024140025	1880	10-15 Floors
408695	1890	13.0	1002127503	1890	10-15 Floors
377755	1890	13.0	1008320066	1890	10-15 Floors

```
[214]: fig = go.Figure(go.Histogram2d(
    x = ny2['decadebuilt'],
    y = ny2['level'],
    nbinsx = 20))
```

```
#fig.update_layout(yaxis_type="log")
fig.show()
```

Removed floors less than 10 to show unusually tall buildings

```
[91]: yearbins = 200
      floorbins = 200

      yearBuiltCut = pd.cut(ny['YearBuilt'], np.linspace(ny['YearBuilt'].min(),
      ↪ny['YearBuilt'].max(), yearbins))
      numFloorsCut = pd.cut(ny['NumFloors'], np.logspace(1, np.log(ny['NumFloors'].
      ↪max()), floorbins))


      xlabel = np.floor(np.linspace(ny['YearBuilt'].min(), ny['YearBuilt'].max(),
      ↪yearbins))
      ylabel = np.floor(np.logspace(1, np.log(ny['NumFloors'].max()), floorbins))

      data = [
          go.Heatmap(z = ny.groupby([numFloorsCut, yearBuiltCut])['BBL'].count().
      ↪unstack().fillna(0).values,
                      colorscale = 'Greens', x = xlabel, y = ylabel)
      ]

      py.iplot(data)
```

```
[118]: cvs = ds.Canvas(800, 500, x_range = (ny['YearBuilt'].min(), ny['YearBuilt'].
      ↪max()),
                      y_range = (ny['NumFloors'].min(),
      ↪ny['NumFloors'].max()))
      agg = cvs.points(ny, 'YearBuilt', 'NumFloors')
      view = tf.shade(agg, cmap = cm.Greys9, how='log')
      export(tf.spread(view, px=2), 'yearvsnumfloors')
```


[118]:



[]:

```
[130]: NewYorkCity = (( -74.29, -73.69), (40.49, 40.92))
cvs = ds.Canvas(700, 700)
agg = cvs.points(ny, 'XCoord', 'YCoord')
view = tf.shade(agg, cmap = cm(inferno))
export(tf.spread(view, px=2), 'firery')
```

[130]:



[133]:


Question You work for a real estate developer and are researching underbuilt areas of the city. After looking in the Pluto data dictionary, you've discovered that all tax assessments consist of two parts: The assessment of the land and assessment of the structure. You reason that there should be a correlation between these two values: more valuable land will have more valuable structures on them (more valuable in this case refers not just to a mansion vs a bungalow, but an apartment tower vs a single family home). Deviations from the norm could represent underbuilt or overbuilt areas of the city. You also recently read a really cool blog post about bivariate choropleth maps, and think the technique could be used for this problem.

Datashader is really cool, but it's not that great at labeling your

visualization. Don't worry about providing a legend, but provide a quick explanation as to which areas of the city are overbuilt, which areas are underbuilt, and which areas are built in a way that's properly correlated with their land value.

```
[139]: NewYorkCity = (( -74.29, -73.69), (40.49, 40.92))
cvs = ds.Canvas(700, 700)
agg = cvs.points(ny, 'XCoord', 'YCoord', ds.mean('AssessLand'))
view = tf.shade(agg, cmap = cm.inferno)
export(tf.spread(view, px=2), 'firery')
```


[139]:



```
[137]: NewYorkCity = (( -74.29, -73.69), (40.49, 40.92))
cvs = ds.Canvas(700, 700)
```

```
agg = cvs.points(ny, 'XCoord', 'YCoord', ds.count('BBL'))
view = tf.shade(agg, cmap = cm(inferno))
export(tf.spread(view, px=2), 'firery')
```

[137]:




```
[218]: ny['assessbldg'] = ny['AssessTot']-ny['AssessLand']
ny['assessRatio']=ny['assessbldg']/ny['AssessTot']
```

Answer Created a new column with ratio of building assessment to total assessment
The brighter areas on the map have a higher ratio meaning that the buildings are overvalued compared to the land (and vice versa)

```
[222]: NewYorkCity = (( -74.29, -73.69), (40.49, 40.92))
cvs = ds.Canvas(700, 700)
```

```
agg = cvs.points(ny, 'XCoord', 'YCoord', ds.count('assessRatio'))  
view = tf.shade(agg, cmap = cm(inferno))  
export(tf.spread(view, px=2), 'firery')
```

[222]:



[]:

[]: